

Complex Data - lab3


Stanisław Wilczyński

Data input and Calculating Means

```
wtloss <- read.table("../data/weightloss.dat",  
header=F)  
## Give names to variables  
names(wtloss) <- c("id", paste("y", 1:4, sep=""), "program")  
## Univariate format  
wtloss.uni <- data.frame(id=rep(wtloss$id, each=4),  
wtloss=as.numeric(t(as.matrix(wtloss[,2:5]))),  
program=rep(wtloss$program, each=4),  
month=seq(0,9,3),  
time.cat=rep(1:4))  
wtloss.uni$prog.fac <- factor(wtloss.uni$program, labels=c("1:encourage", "2:none"))  
attach(wtloss.uni)  
wgt.mean <- tapply(wgtloss, list(month, program), mean)  
wgt.sd <- tapply(wgtloss, list(month, program), sd)  
detach(wtloss.uni)
```

Plotting Response Profiles

```
xyplot(wgtloss~month|prog.fac, type='l', groups=id, data=wtloss.uni)
```



Choice of covariance structure

We will consider three covariance structures:

1. Unstructured (UN)
2. Compound symmetry (CS)
3. Autoregressive (AR-1)

```
## CATEGORICAL time, UN, REML
wtloss.un.cat <- gls(wgtloss~factor(month)*prog.fac,
correlation=corSymm(form= ~1 | id),
weights=varIdent(form= ~1 | month),
data=wtloss.uni)
#summary(wtloss.un.cat)

## CATEGORICAL time, CS, REML
wtloss.cs.cat <- gls(wgtloss~factor(month)*prog.fac,
correlation=corCompSymm(form= ~1 | id),
weights=varIdent(form= ~1),
data=wtloss.uni)
#summary(wtloss.cs.cat)

## CATEGORICAL time, AR1, REML
wtloss.ar1.cat <- gls(wgtloss~factor(month)*prog.fac,
correlation=corAR1(form= ~1 | id),
weights=varIdent(form= ~1),
data=wtloss.uni)
#summary(wtloss.ar1.cat)
```

Compound symmetry vs unstructured

H_0 : compound symmetry is adequate for the data

H_1 : unstructured is required

```
un_vs_cs <- anova(wtloss.un.cat, wtloss.cs.cat)
```

Structure	CS vs. UN	
	-2 REML Log-Likelihood	Number of Cov. Parameters
Compound Symmetry	1905.8232573	10
Unstructured	1900.8766893	18
Difference	4.9465681	8

LRT yields $G^2 = 4.9465681$ with 8 df (p-value = 0.7632715) so we do not reject the null hypothesis at $\alpha = 0.05$ and conclude that the assumption of compound symmetry covariance structure is adequate for the data.

Autoregressive vs unstructured

H_0 : autoregressive is adequate for the data

H_1 : unstructured is required

```
un_vs_ar <- anova(wtloss.un.cat, wtloss.ar1.cat)
```

AR-1 vs. UN		
	-2 REML	Number of
Structure	Log-Likelihood	Cov. Parameters
Autoregressive	1931.246837	10
Unstructured	1900.8766893	18
Difference	30.3701477	8

LRT yields $G^2 = 30.3701477$ with 8 df (p-value = 1.8178034×10^{-4}), so we reject the null hypothesis at $\alpha = 0.05$ and conclude that the assumption of autoregressive covariance structure is not adequate when compared to unstructured.

Compound symmetry vs autoregressive

```
cs_vs_ar1 <- anova(wtloss.cs.cat, wtloss.ar1.cat)
```

Since CS and AR-1 have the same number of parameters = 10, no LRT is necessary. We can directly compare their likelihoods, or $-2 \log(\text{likelihood})$:

- $-2 * \log(\text{likelihood})$ for CS = 1905.8232573
- $-2 * \log(\text{likelihood})$ for AR-1 = 1931.246837

Since $-2 * \log(\text{likelihood})$ for CS is smaller than for AR-1, CS has a higher likelihood and we conclude that CS is an adequate model for the covariance structure when compared to AR-1.

The most adequate covariance structure is compound symmetry model.

Are all these tests correct?

No, the last test is theoretically not correct, because for non nested models we shouldn't draw conclusions just basing on comparison of log-likelihoods or performing LRT. We should compare AIC or BIC instead. However, in this case the number of parameters is the same, so comparing likelihoods is equivalent to comparing AICs. Moreover, when comparing many models instead of pairwise tests we can use a better strategy and choose the one with the lowest value of AIC.

	-2 REML	Number of	
Structure	Log-Likelihood	Parameters	AIC
Compound Symmetry	1905.8232573	10	1925.8232573
Autoregressive	1931.246837	10	1951.246837
Unstructured	1900.8766893	18	1936.8766893

Thus, we will use a compound symmetry covariance structure for the remainder of the lab.

Single Degree of Freedom Contrasts

AUC - test for equality of the area under the curve in two groups.

```
t <- wtloss.uni$month
L2 <- 0.5*c(t[1] + t[2] - 2*t[4], t[3] - t[1], t[4]-t[2], t[4]-t[3])
coefs <- summary(wtloss.cs.cat)$coefficients
#Estimated mean AUC in encouragement program
AUC_enc <- sum(L2 * wgt.mean[,1])
#Estimated mean AUC in no encouragement program
AUC_none <- sum(L2 * wgt.mean[,2])
```

The obtained values of AUCs are:

Estimated mean AUC in encouragement program is -319.5705882.

Estimated mean AUC in no encouragement program is -117.7038462.

As written in the textbook to test for the equality of AUCs, we employ the contrast $(-L_2, L_2)$.

```
L <- c(-L2, L2)
anova(wtloss.cs.cat, L=L)

## Denom. DF: 232
## F-test for linear combination(s)
##              (Intercept)              factor(month)3
##              7.5              -3.0
##              factor(month)6              factor(month)9
##              -3.0              -1.5
##              prog.fac2:none factor(month)3:prog.fac2:none
##              -7.5              3.0
## factor(month)6:prog.fac2:none factor(month)9:prog.fac2:none
##              3.0              1.5
##   numDF  F-value p-value
## 1      1 759.3707 <.0001
```

The p-value is almost 0, so we reject the null hypothesis that the response profile is the same for two treatments.

Parametric curves

Quadratic time trend

```
## QUADRATIC time, CS, ML,
wtloss.cs.quad <- gls(wgtloss~month*prog.fac + I(month^2)*prog.fac,
correlation=corCompSymm(form= ~1 | id),
weights=varIdent(form= ~1),
data=wtloss.uni,
method="ML")

## CATEGORICAL time, CS, ML
wtloss.cs.cat.ml <- gls(wgtloss~factor(month)*prog.fac,
correlation=corCompSymm(form= ~1 | id),
weights=varIdent(form= ~1),
data=wtloss.uni,
method="ML")
```

H_0 : quadratic model

H_1 : saturated mode

```
qm_vs_sm <- anova(wtloss.cs.cat.ml, wtloss.cs.quad)
```

Testing the Quadratic trend		
Structure	-2 Log Likelihood	Number of Parameters
Quadratic model	1935.7000029	8
Saturated model	1935.3313128	10
Difference	0.3686901	2

LRT yields $G^2 = 0.3686901$ with 2 df (p-value = 0.8316488), so we fail to reject the null hypothesis at $\alpha = 0.05$ and conclude that the model with Month as a quadratic effect seems to fit the data adequately. (Note: $\chi^2_{2,0.95} = 5.99$)

Linear time trend

```
wtloss.cs.lin <- gls(wgtloss~month*prog.fac,
correlation=corCompSymm(form= ~1 | id),
weights=varIdent(form= ~1),
data=wtloss.uni,
method="ML")
```

H_0 : linear model

H_1 : quadratic model

```
qm_vs_lm <- anova(wtloss.cs.quad, wtloss.cs.lin)
```

Linear vs. Quadratic		
Structure	-2 Log Likelihood	Number of Parameters
Linear model	1936.295409	6
Quadratic model	1935.7000029	8
Difference	0.5954061	2

LRT yields $G^2 = 0.5954061$ with 2 df (p-value = 0.7425218), so we fail to reject the null hypothesis at $\alpha = 0.05$ and conclude that the model with Month as a linear effect fits the data better than the model with Month as a quadratic effect. (Note: $\chi^2_{2,0.95} = 5.99$)

Testing for intersections in the linear model

```
## Linear model, NO interactions
wtloss.cs.lin.noint <- gls(wgtloss~month + prog.fac,
correlation=corCompSymm(form= ~1 | id),
weights=varIdent(form= ~1),
data=wtloss.uni,
method="ML")
```

$H_0 : \beta_4 = 0$

$H_1 : \beta_4 \neq 0$

```
anova(wtloss.cs.lin)
```

```
## Denom. DF: 236
##          numDF  F-value p-value
## (Intercept)      1 6970.808 <.0001
## month           1 1011.995 <.0001
## prog.fac        1   24.301 <.0001
## month:prog.fac   1  184.368 <.0001
```

#we employ proper contrast

```
anova(wtloss.cs.lin, L=c(0,0,0,1))
```

```
## Denom. DF: 236
```

```
## F-test for linear combination(s)
```

```
## [1] 1
```

```
## numDF F-value p-value
## 1      1 184.3676 <.0001
```

Since the p-value for month*program is almost 0, we reject the null hypothesis and conclude that there is a significant interaction between month and program.

Thus, our final model is the linear model with interaction.

```
final.wtloss.cs.lin <- gls(wgtloss~month*prog.fac,
correlation=corCompSymm(form= ~1 | id),
weights=varIdent(form= ~1),
data=wtloss.uni)
summary(final.wtloss.cs.lin)

## Generalized least squares fit by REML
## Model: wgtloss ~ month * prog.fac
## Data: wtloss.uni
##      AIC      BIC    logLik
## 1941.074 1961.857 -964.5368
##
## Correlation Structure: Compound symmetry
## Formula: ~1 | id
## Parameter estimate(s):
##      Rho
## 0.8083168
##
## Coefficients:
##              Value Std.Error   t-value p-value
## (Intercept)    242.37676   3.655475   66.30513  0.0000
## month          -8.10431   0.245748  -32.97814  0.0000
## prog.fac2:none    2.99477   5.553065    0.53930  0.5902
## month:prog.fac2:none  5.08329   0.373318   13.61650  0.0000
##
## Correlation:
##              (Intr) month  prg.2:
## month          -0.303
## prog.fac2:none  -0.658  0.199
## month:prog.fac2:none  0.199 -0.658 -0.303
##
## Standardized residuals:
##      Min      Q1      Med      Q3      Max
## -3.0172539 -0.6942911 -0.0369479  0.7699925  2.2892673
##
## Residual standard error: 21.95552
## Degrees of freedom: 240 total; 236 residual
```

In terms of effectiveness we can conclude that the first program is much more effective than the second one: in the first group the estimated rate of change equals $\hat{\beta}_2 = -8.1043137$ whereas in the second group estimated rate of change is $\hat{\beta}_2 + \hat{\beta}_4 = -3.0210256$.

Interpreting quadratic trends:

```
wtloss.cs.quad.noint <- gls(wgtloss-month + prog.fac + I(month^2),
correlation=corCompSymm(form= ~1 | id),
weights=varIdent(form= ~1),
data=wtloss.uni)
ncoefs <- wtloss.cs.quad.noint$coefficients
coefs <- c(ncoefs[1], ncoefs[2], ncoefs[4], ncoefs[3])
change_rates <- coefs[2] + 2*coefs[3]*t[1:4]
expected1 <- coefs[1] + coefs[2]*t[1:4] + coefs[3]*t[1:4]*t[1:4]
```

$\hat{\beta}_1 = 232.0118529$ expected mean response at baseline of subjects in program 1.

$\hat{\beta}_4 = 25.8695701$ change in expected response at baseline of subjects in program 2 vs. program 1.

Rate of change in program 1 is $\beta_2 + 2\beta_3 time_{ij}$.

Plugging in the above estimates,

Program 1		
time	Rate of Change	Expected Response
0	-5.4490556	232.0118529
1	-5.7507222	215.2121863
2	-6.0523889	197.5075196
3	-6.3540556	178.8978529

Thus, the mean response for Program 1 decreases over time. For the second program the effect is exactly the same - in this case for every time point 25.8695701 is added to the expected response, because in our model we don't include any interactions and the group effect is incorporated just by adding a constant.